

VE 492 Homework8

Due: 23:59, July 28

Question 1: Maximum Likelihood Estimation

We will begin with a short derivation. Consider a probability distribution with a domain that consists of $|X|$ different values. We get to observe N total samples from this distribution. We use n_i to represent the number of the N samples for which outcome i occurs. Our goal is to estimate the probabilities $\theta_i, i = 1, 2, \dots, |X| - 1$ of each of the events. The probability of the last outcome, $|X|$, equals $1 - \sum_{i=1}^{|X|-1} \theta_i$.

In *maximum likelihood estimation*, we choose the θ_i that maximize the likelihood of the observed samples,

$$L(\text{samples}, \theta) \propto (1 - \theta_1 - \theta_2 - \dots - \theta_{|X|-1})^{n_{|X|}} \prod_{i=1}^{|X|-1} \theta_i^{n_i}$$

For this derivation, it is easiest to work with the log of the likelihood. Maximizing log-likelihood also maximizes likelihood, since the quantities are related by a monotonic transformation. Taking logs we obtain

$$\theta^{\text{ML}} = \underset{\theta}{\operatorname{argmax}} n_{|X|} \log(1 - \theta_1 - \theta_2 - \dots - \theta_{|X|-1}) + \sum_{i=1}^{|X|-1} n_i \log \theta_i$$

Setting derivatives with respect to θ_i equal to zero, we obtain $|X| - 1$ equations in the $|X| - 1$ unknowns, $\theta_1, \theta_2, \dots, \theta_{|X|-1}$:

$$\frac{-n_{|X|}}{1 - \theta_1^{\text{ML}} - \theta_2^{\text{ML}} - \dots - \theta_{|X|-1}^{\text{ML}}} + \frac{n_i}{\theta_i^{\text{ML}}} = 0$$

Multiplying by $\theta_i(1 - \theta_1 - \theta_2 - \dots - \theta_{|X|-1})$ makes the original $|X| - 1$ nonlinear equations into $|X| - 1$ linear equations:

$$-n_{|X|} \theta_i^{\text{ML}} + n_i (1 - \theta_1^{\text{ML}} - \theta_2^{\text{ML}} - \dots - \theta_{|X|-1}^{\text{ML}}) = 0$$

That is, the maximum likelihood estimation of θ can be found by solving a linear system of $|X| - 1$ equations in $|X| - 1$ unknowns. Doing so shows that the maximum likelihood estimate corresponds to simply the count for each outcome divided by the total number of samples. I.e., we have that:

$$\theta_i^{\text{ML}} = \frac{n_i}{N}$$

Notice: Please write each sub-question in one row, that is, there will be 3 rows for this question. And please use irreducible fractions for your answer.

Sample Answer:

1, 1/2, 1/3, 1/4

2/5 (instead of 4/10), 1/3, 4/7

3/8, 3/7, 3/5

Part 1.

Now, consider a sampling process with 3 possible outcomes: R, G, and B. We observe the following sample counts:

outcome	R	G	B
count	3	1	7

- 1) What is the total sample count N ?
- 2) What are the maximum likelihood estimates for the probabilities of each outcome?

$$\theta_R^{ML} =$$

$$\theta_G^{ML} =$$

$$\theta_B^{ML} =$$

Part 2.

Now, use *Laplace smoothing* with strength $k = 3$ to estimate the probabilities of each outcome.

$$\theta_R^{LAP,3} =$$

$$\theta_G^{LAP,3} =$$

$$\theta_B^{LAP,3} =$$

Part 3.

Now, consider Laplace smoothing in the limit $k \rightarrow \infty$. Fill in the corresponding probability estimates.

$$\theta_R^{LAP,\infty} =$$

$$\theta_G^{LAP,\infty} =$$


$$\theta_B^{LAP,\infty} =$$

Question 2: Poisson Parameter Evaluation

We will now consider maximum likelihood estimation in the context of a different probability distribution. Under the Poisson distribution, the probability of an event occurring $X = k$ times is:

$$P(X = k) = \frac{v^k e^{-v}}{k!}$$

Here v is the parameter we wish to estimate. The distribution is plotted for several values of v below.



On a sheet of scratch paper, work out the maximum likelihood estimate for v , given observations of several k_i .


Hints: start by taking the product of the equation above over all the k_i , and then taking the log. Then, differentiate with respect to v , set the result equal to 0, and solve for v in terms of the k_i .

You observe the samples $k_1 = 6, k_2 = 3, k_3 = 8, k_4 = 4, k_5 = 2$. What is your maximum likelihood estimate of v ?

Sample Answer (rounded to 3 decimal places):
0.160

Question 3: Naive Bayes

In this question, we will train a Naive Bayes classifier to predict class labels Y as a function of input features F_i .



We are given the following 15 training points:

F_1	1	1	1	1	1	1	1	0	1	1	1	0	1	1	0
F_2	0	0	0	1	0	0	0	0	0	1	0	0	0	0	1
F_3	0	0	0	0	0	0	0	0	0	0	1	0	1	0	0
Y	A	A	A	A	A	A	A	A	A	A	B	B	B	B	C

Note: Please write your answer for each table in one row, that is, there will be 10 rows for this question. Besides, please use values rounded to 3 decimal places.

Sample Answer:

0.160,0.170,0.160

...

...

...

A

...

0.200,0.211,...

1) What is the maximum likelihood estimate of the prior $P(Y)$?

Y	$P(Y)$
A	
B	
C	

2) What are the maximum likelihood estimates of the conditional probability distributions? Fill in the tables below (the second and third are done for you).

F_1	Y	$P(F_1 Y)$
0	A	
1	A	
0	B	
1	B	
0	C	
1	C	

F_2	Y	$P(F_2 Y)$
0	A	0.800
1	A	0.200
0	B	1.000
1	B	0.000
0	C	0.000
1	C	1.000

F_3	Y	$P(F_3 Y)$
0	A	1.000
1	A	0.000
0	B	0.500
1	B	0.500
0	C	1.000
1	C	0.000

- 3) Now consider a new data point ($F_1 = 0, F_2 = 0, F_3 = 1$). Use your classifier to determine the joint probability of causes Y and this new data point, along with the posterior probability of Y given the new data:

Y	$P(Y, F_1 = 0, F_2 = 0, F_3 = 1)$
A	
B	
C	

Y	$P(Y F_1 = 0, F_2 = 0, F_3 = 1)$
A	
B	
C	

- 4) What label does your classifier give to the new data point? (Break ties alphabetically). Write capital letters only.
- 5) Now use Laplace Smoothing with strength $k = 3$ to estimate the prior $P(Y)$ for the same data.

Y	$P(Y)$
A	0.542
B	0.292
C	0.167

$$(10+3)/(15+3*3)$$

- 6) Use Laplace Smoothing with strength $k = 3$ to estimate the conditional probability distributions below (again, the second two are done for you).

$$(1+3)/(10+2*3)$$

F_1	Y	$P(F_1 Y)$
0	A	0.250
1	A	0.750
0	B	0.400
1	B	0.600
0	C	0.571
1	C	0.429

F_2	Y	$P(F_2 Y)$
0	A	0.688
1	A	0.312
0	B	0.700
1	B	0.300
0	C	0.429
1	C	0.571

F_3	Y	$P(F_3 Y)$
0	A	0.812
1	A	0.188
0	B	0.500
1	B	0.500
0	C	0.571
1	C	0.429

- 7) Now consider again the new data point ($F_1 = 0, F_2 = 0, F_3 = 1$). Use the Laplace-Smoothed version of your classifier to determine the joint probability of causes Y and this new data point, along with the posterior probability of Y given the new data:

Y	$P(Y, F_1 = 0, F_2 = 0, F_3 = 1)$
A	0.018
B	0.041
C	0.018

Y	$P(Y F_1 = 0, F_2 = 0, F_3 = 1)$
A	0.234
B	0.532
C	0.234

- 8) What label does your (Laplace-Smoothed) classifier give to the new data point? (Break ties alphabetically). Write a single capital letter.

Question 4: Datasets

When training a classifier, it is common to split the available data into a training set, a hold-out set, and a test set, each of which has a different role.

Sample Answer:

A


A

A

- 1) Which data set is used to learn the conditional probabilities?
 - A. Training Data
 - B. Hold-Out Data
 - C. Test Data
- 2) Which data set is used to tune the Laplace Smoothing hyperparameters?
 - A. Training Data
 - B. Hold-Out Data
 - C. Test Data
- 3) Which data set is used for quantifying performance results?
 - A. Training Data
 - B. Hold-Out Data
 - C. Test Data

Question 5: Linear Separability

Consider the data in the figure below.



The data is plotted as a function of two features, f_1 and f_2 . As plotted, the data is not linearly separable. Which of the following candidate features f_3 , when added, would cause the data to be linearly separable? Choose all possible answer(s).

- A. $f_3 = |f_1| + |f_2|$
- B. $f_3 = \sin(f_1)$
- C. $f_3 = f_1^2 + f_2^2$
- D. $f_3 = f_1^2$
- E. $f_3 = f_1$
- F. $f_3 = 1$
- G. $f_3 = f_1 f_2$
- H. $f_3 = 1$ if $f_1 \in [-7, 7]$ and $f_2 \in [-7, 7]$, 0 otherwise